



Functional Performance

Martin Thompson - @mjpt777

**Are all memory
operations equal?**

Sequential Access

-

**Average time in ns/op to sum all
longs in a 1GB array?**

Access Pattern Benchmark


Benchmark	Score	Error	Units
=====			
sequential	0.832	± 0.006	ns/op

~1 ns/op

Really???

Less than 1ns per operation?

Haswell Execution Unit Overview



Access Pattern Benchmark

Benchmark	Score	Error	Units
=====			
sequential	0.832	\pm 0.006	ns/op
randomPage	2.703	\pm 0.025	ns/op
dependentRandomPage	7.102	\pm 0.326	ns/op
randomHeap	19.896	\pm 3.110	ns/op
dependentRandomHeap	89.516	\pm 4.573	ns/op

Access Pattern Benchmark

Benchmark	Score	Error	Units
=====			
sequential	0.832	± 0.006	ns/op
randomPage	2.703	± 0.025	ns/op
dependentRandomPage	7.102	± 0.326	ns/op
randomHeap	19.896	± 3.110	ns/op
dependentRandomHeap	89.516	± 4.573	ns/op

~90 ns/op

Some Fundamental Laws

$$r = s(2 - \rho) / 2(1 - \rho)$$

Queueing Theory

$$r = s (2 - \rho) / 2 (1 - \rho)$$

r = mean response time

s = service time

ρ = utilisation

Queueing Theory

$$r = s (2 - \rho) / 2 (1 - \rho)$$


r = mean response time

s = service time

ρ = utilisation

Note: $\rho = \lambda * s$

Queueing Theory





$$C(N) = N / (1 + \alpha(N - 1) + ((\beta * N) * (N - 1)))$$

Universal Scalability Law (USL)

$$C(N) = N / (1 + \alpha(N - 1) + ((\beta * N) * (N - 1)))$$


C = capacity or throughput

N = number of processors

α = **contention** penalty


β = **coherence** penalty

Universal Scalability Law (USL)





AMD EPYC Infinity Fabric



Source: AMD

Systems Engineering

Messaging

***Concurrent
Distributed***

Network & Protocol Understanding

What are the normal approaches?

Mechanical Sympathy


Skiplists, Trees, Scoreboards

Mechanical Sympathy

Skiplists, Trees, Scoreboards

->

Pointer chasing hell



Immutability ?

Concurrent Tree Updates

\Rightarrow

Path Copy

Garbage Collection Hell



***Functional data structures
are like sausages,
the more you see them being
made, the less you will sleep***

CRDT ???

CRDT ???

Conflict-free Replicated Data Type

CRDT ???

Conflict-free Replicated Data Type

Operation Based

Commutative Replicated Data Type

CmRDT

CRDT ???

Conflict-free Replicated Data Type

Operation Based

Commutative Replicated Data Type

CmRDT

State Based

Convergent Replicated Data Type

CvRDT

Commutative Replicated Data Types

- **Operations are replicated**
 - e.g. Add 10 to value, or -10

Commutative Replicated Data Types

- **Operations are replicated**
 - e.g. Add 10 to value, or -10
- **Operations are commutative**
 - i.e. They can be applied in any order

Commutative Replicated Data Types

- **Operations are replicated**
 - e.g. Add 10 to value, or -10
- **Operations are commutative**
 - i.e. They can be applied in any order
- **Operations are NOT idempotent**
 - Network transport must be reliable

Convergent Replicated Data Types

- **State is replicated**
 - **i.e. Whole data structure is replicated**

Convergent Replicated Data Types


- **State is replicated**
 - **i.e. Whole data structure is replicated**
- **Updates must monotonically increase state**
 - **e.g. Append only, remove requires tombstones**

Convergent Replicated Data Types


- **State is replicated**
 - i.e. Whole data structure is replicated
- **Updates must monotonically increase state**
 - e.g. Append only, remove requires tombstones
- **Merge function must be associative, commutative, and idempotent**
 - A reliable transport is not required
 - Some designs support deltas for efficiency
 - Resolving conflict can be difficult

Aeron

Log Buffer



File

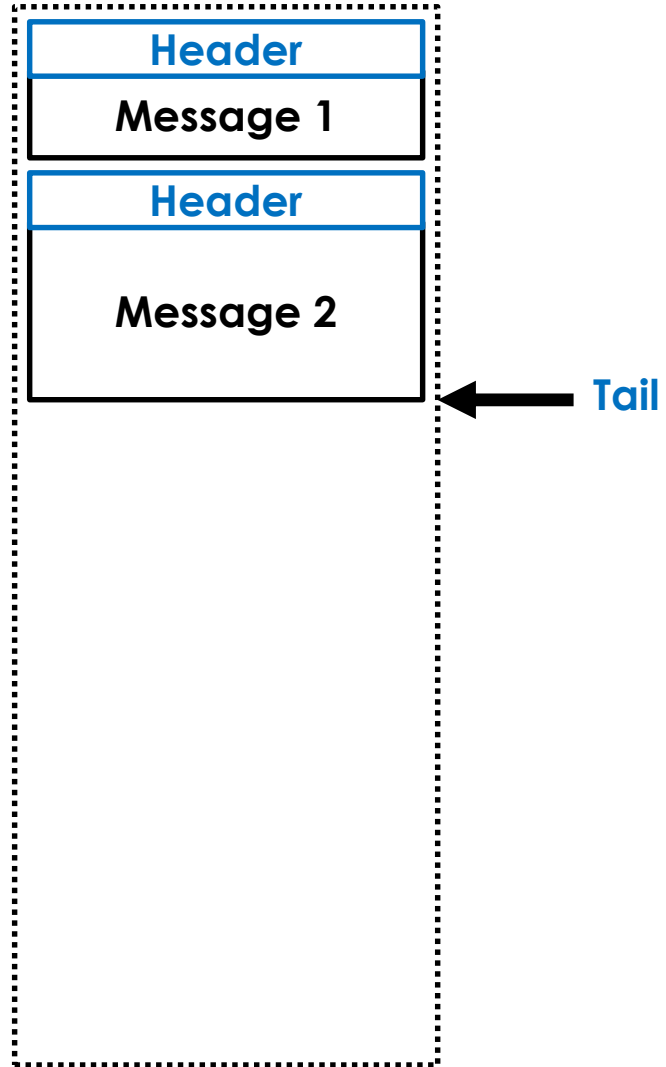


Header


Message 1

Tail

File

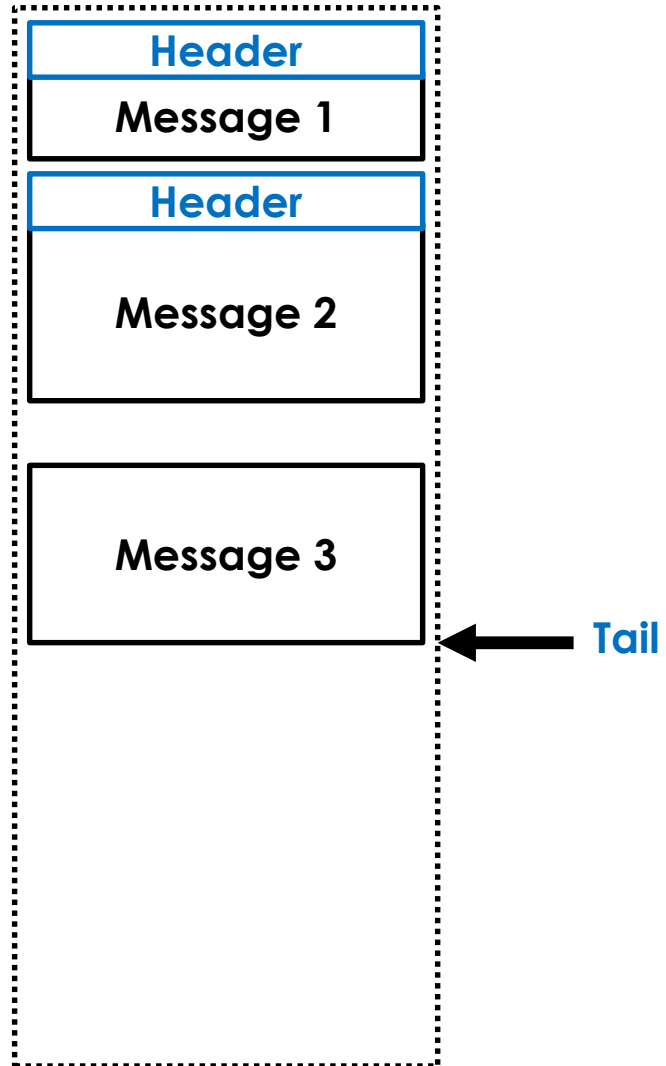


File




← Tail

File



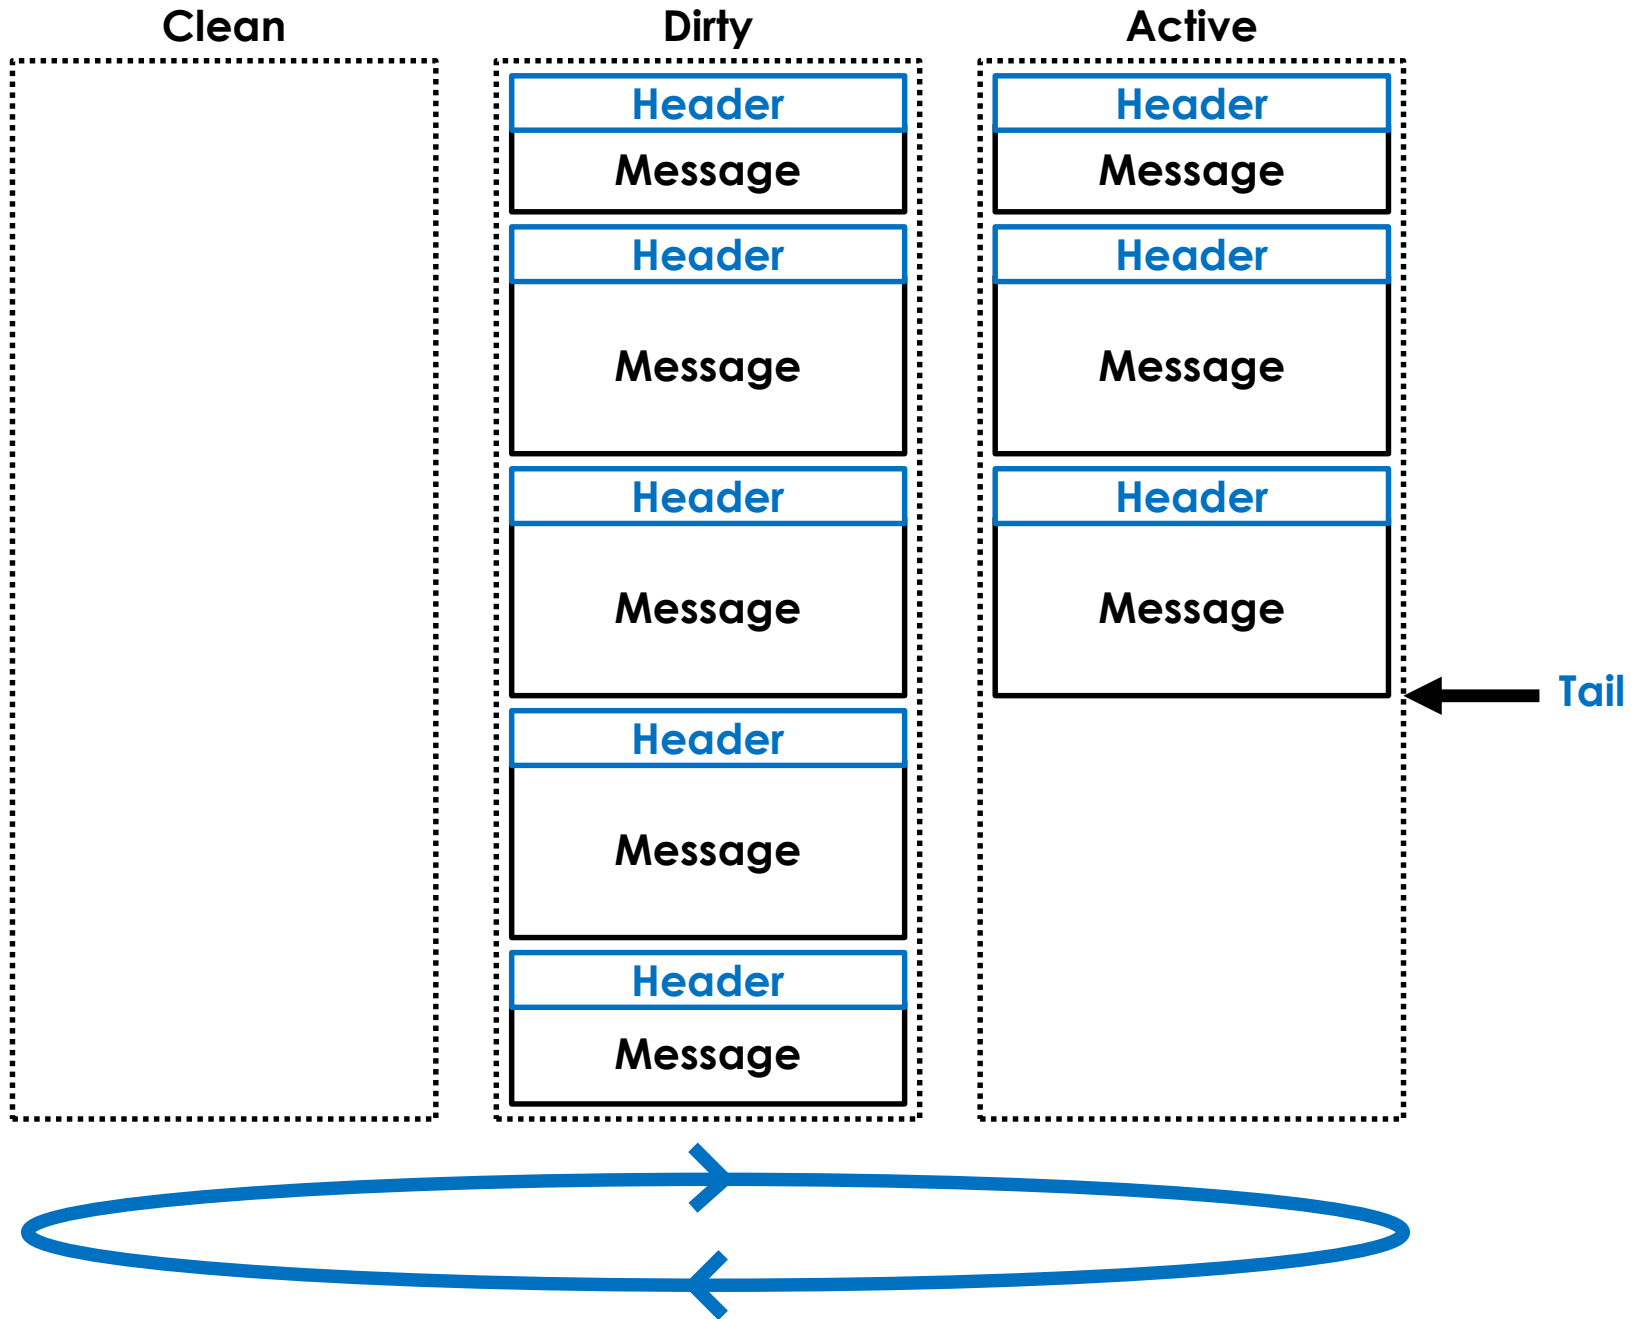
File



***One big file that
goes on forever?***

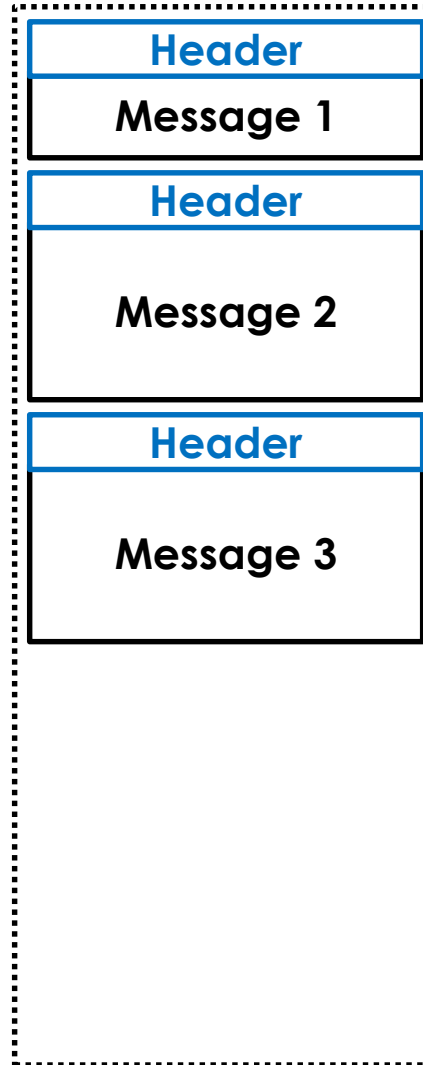
No!!!

***Page faults, page cache churn,
VM pressure, ...***



Concurrent Publication?

File




Tail

Message X

Message Y


File



Message X

Message Y

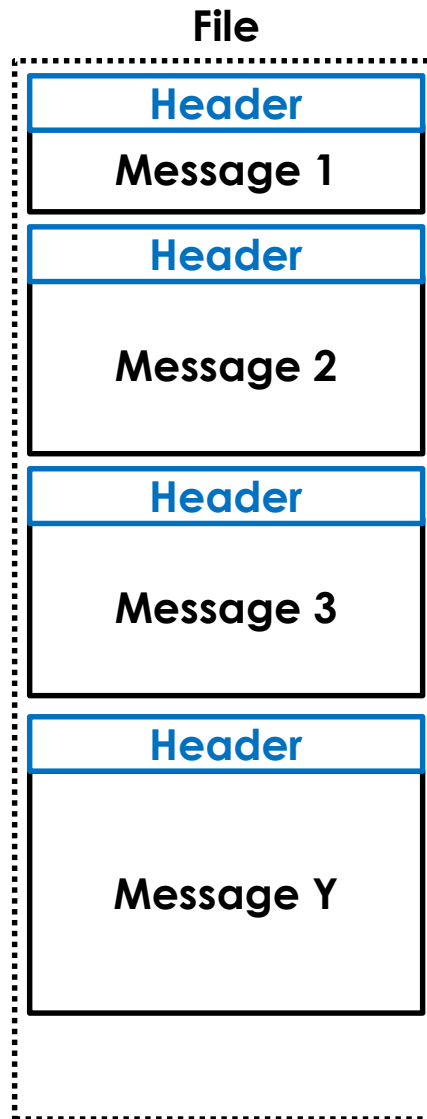
File



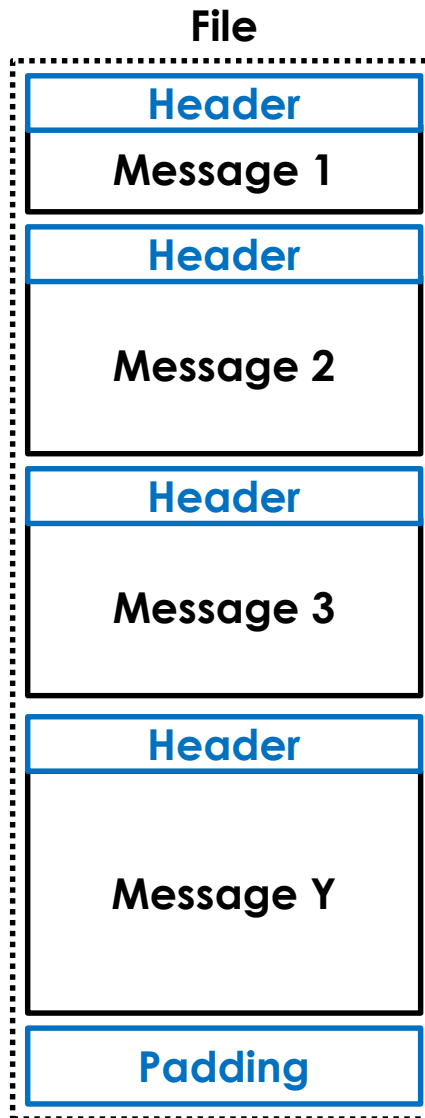
Message X

Message Y



← Tail



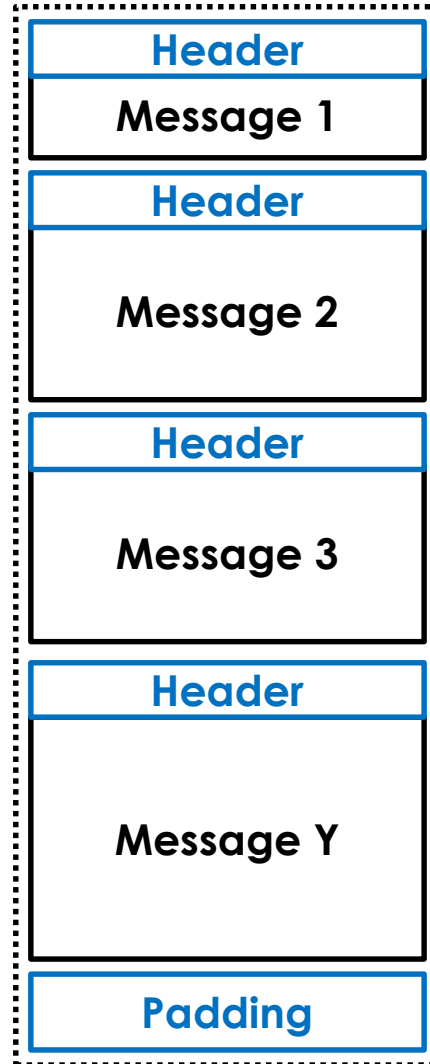
← Tail




← Tail



File

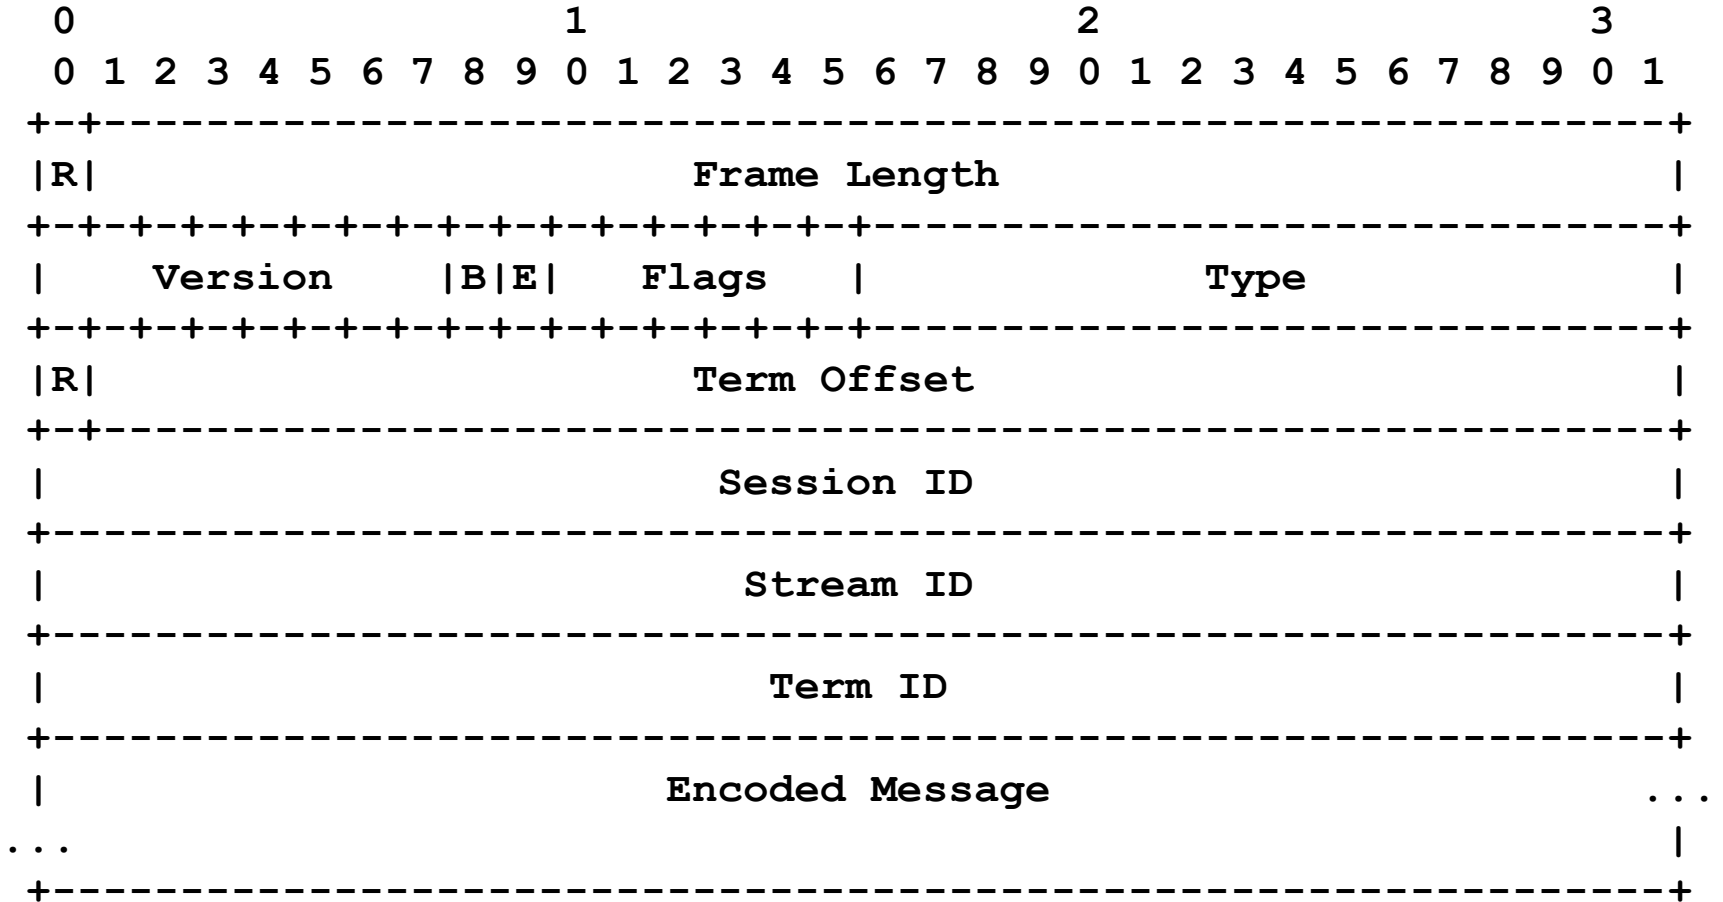


File




What's in a header?

Data Message Header

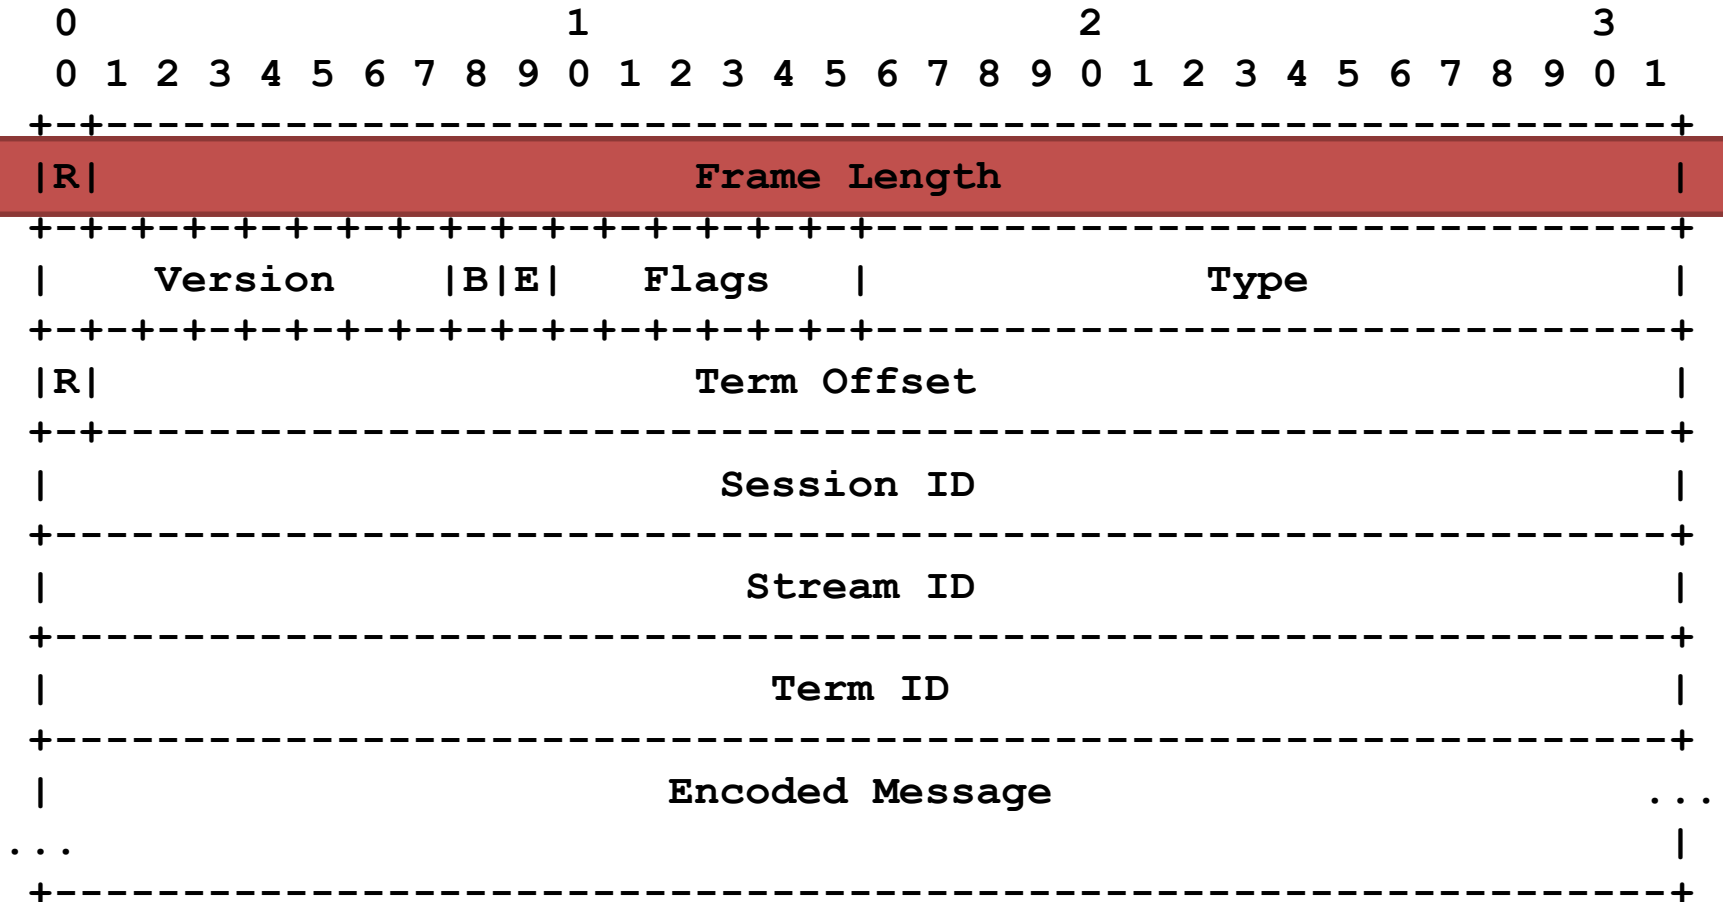


***What if a Publisher dies mid
operation?***

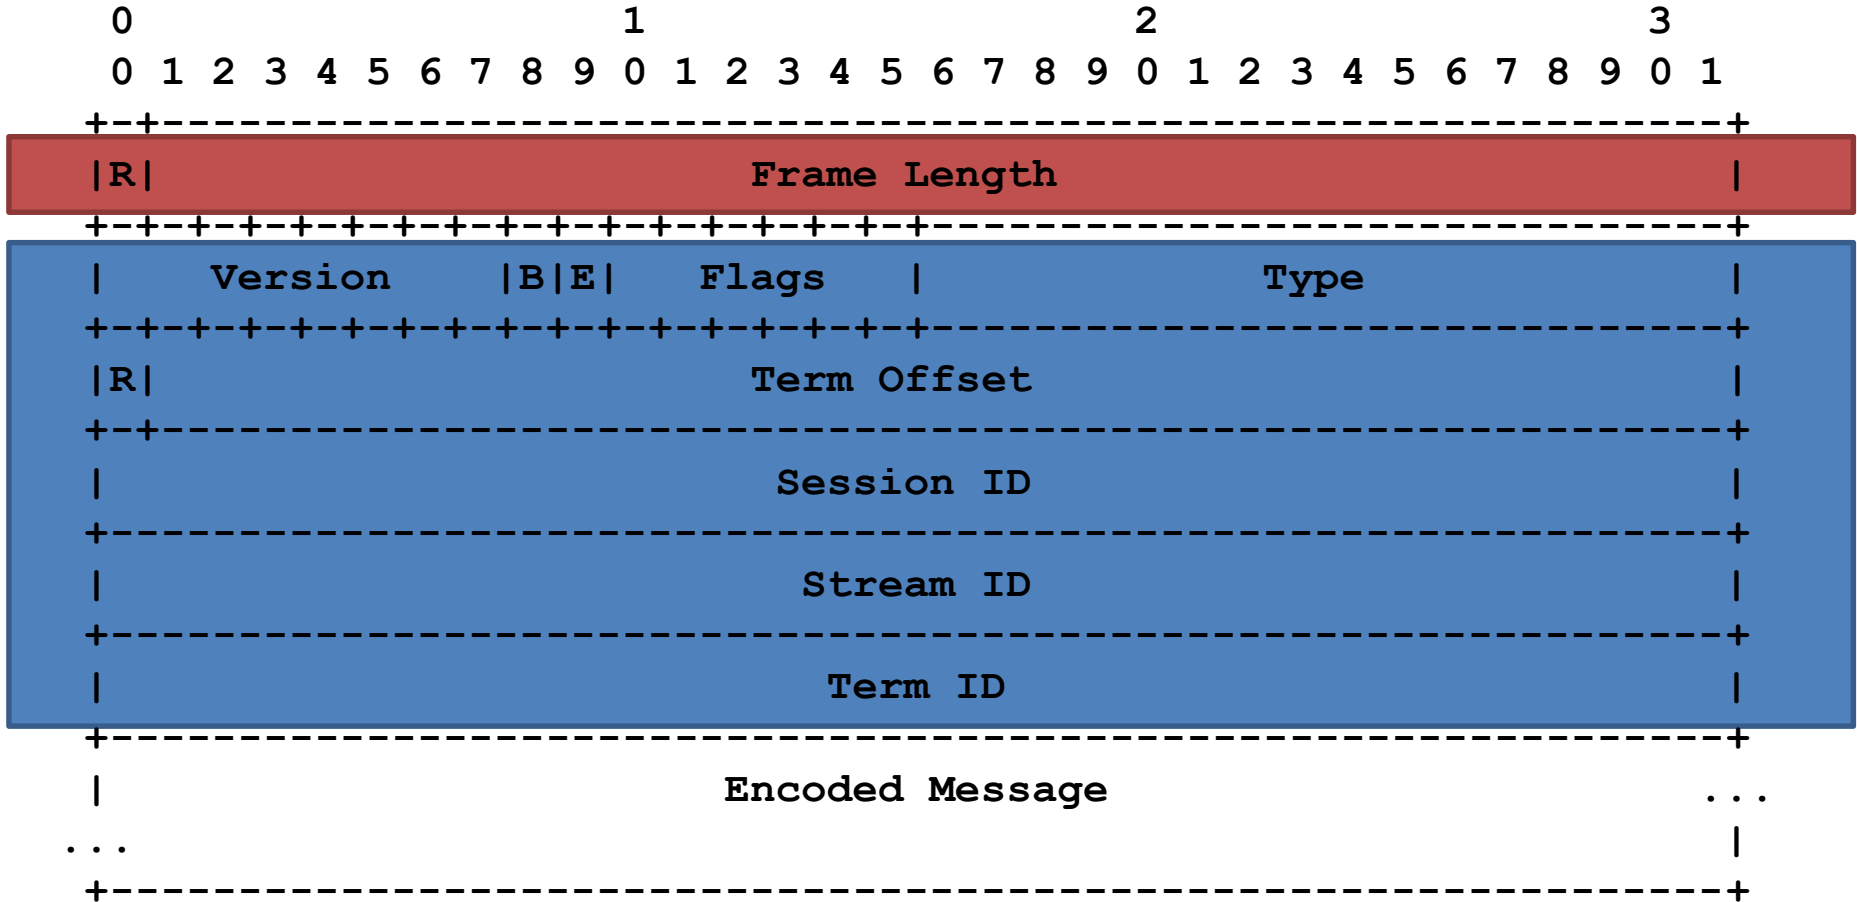
Data Message Header



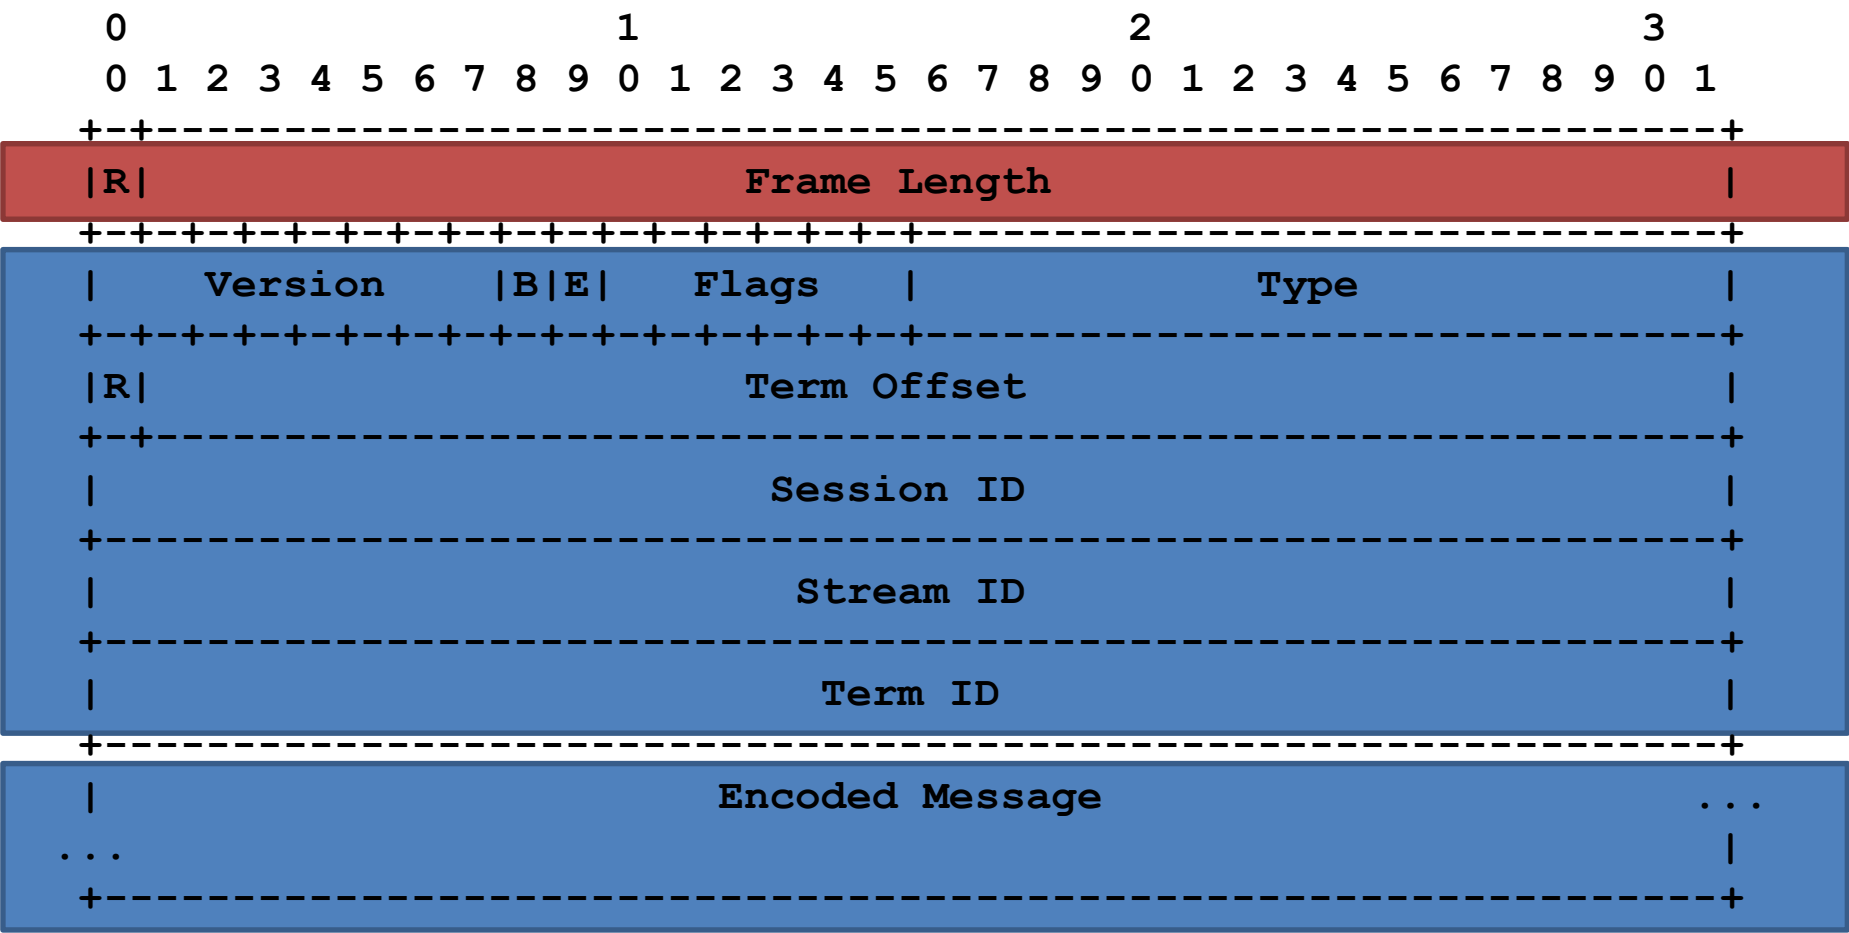
Data Message Header



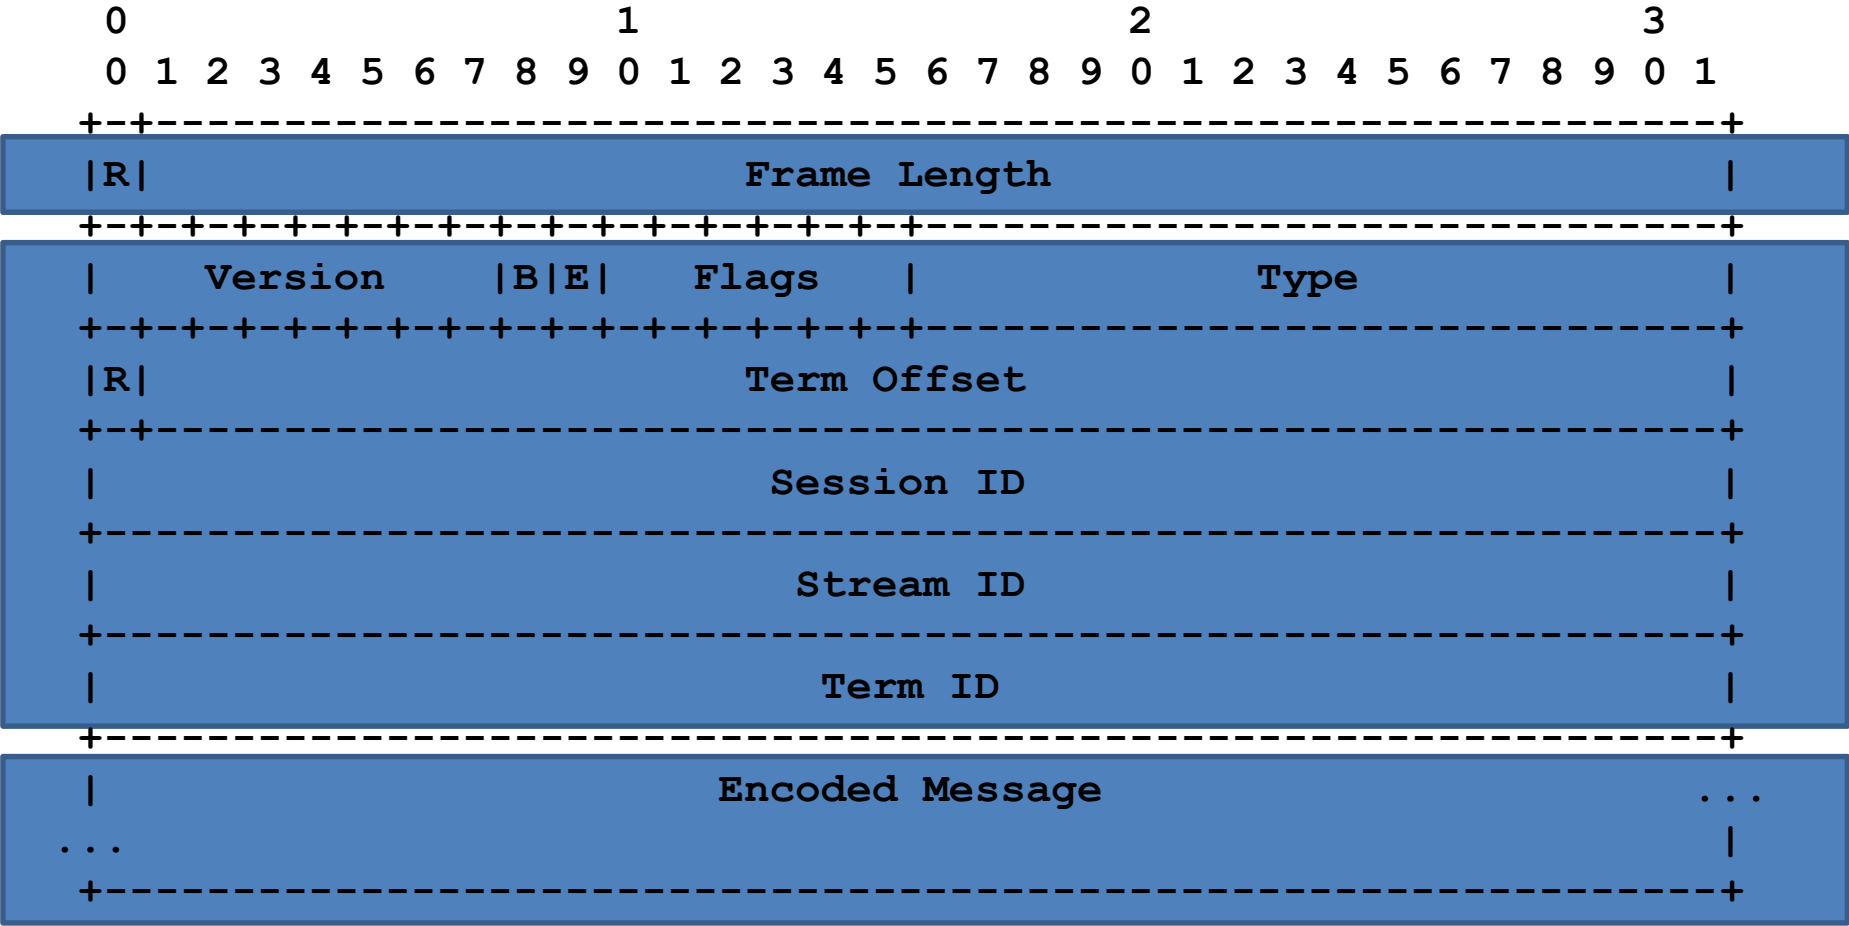
Data Message Header



Data Message Header

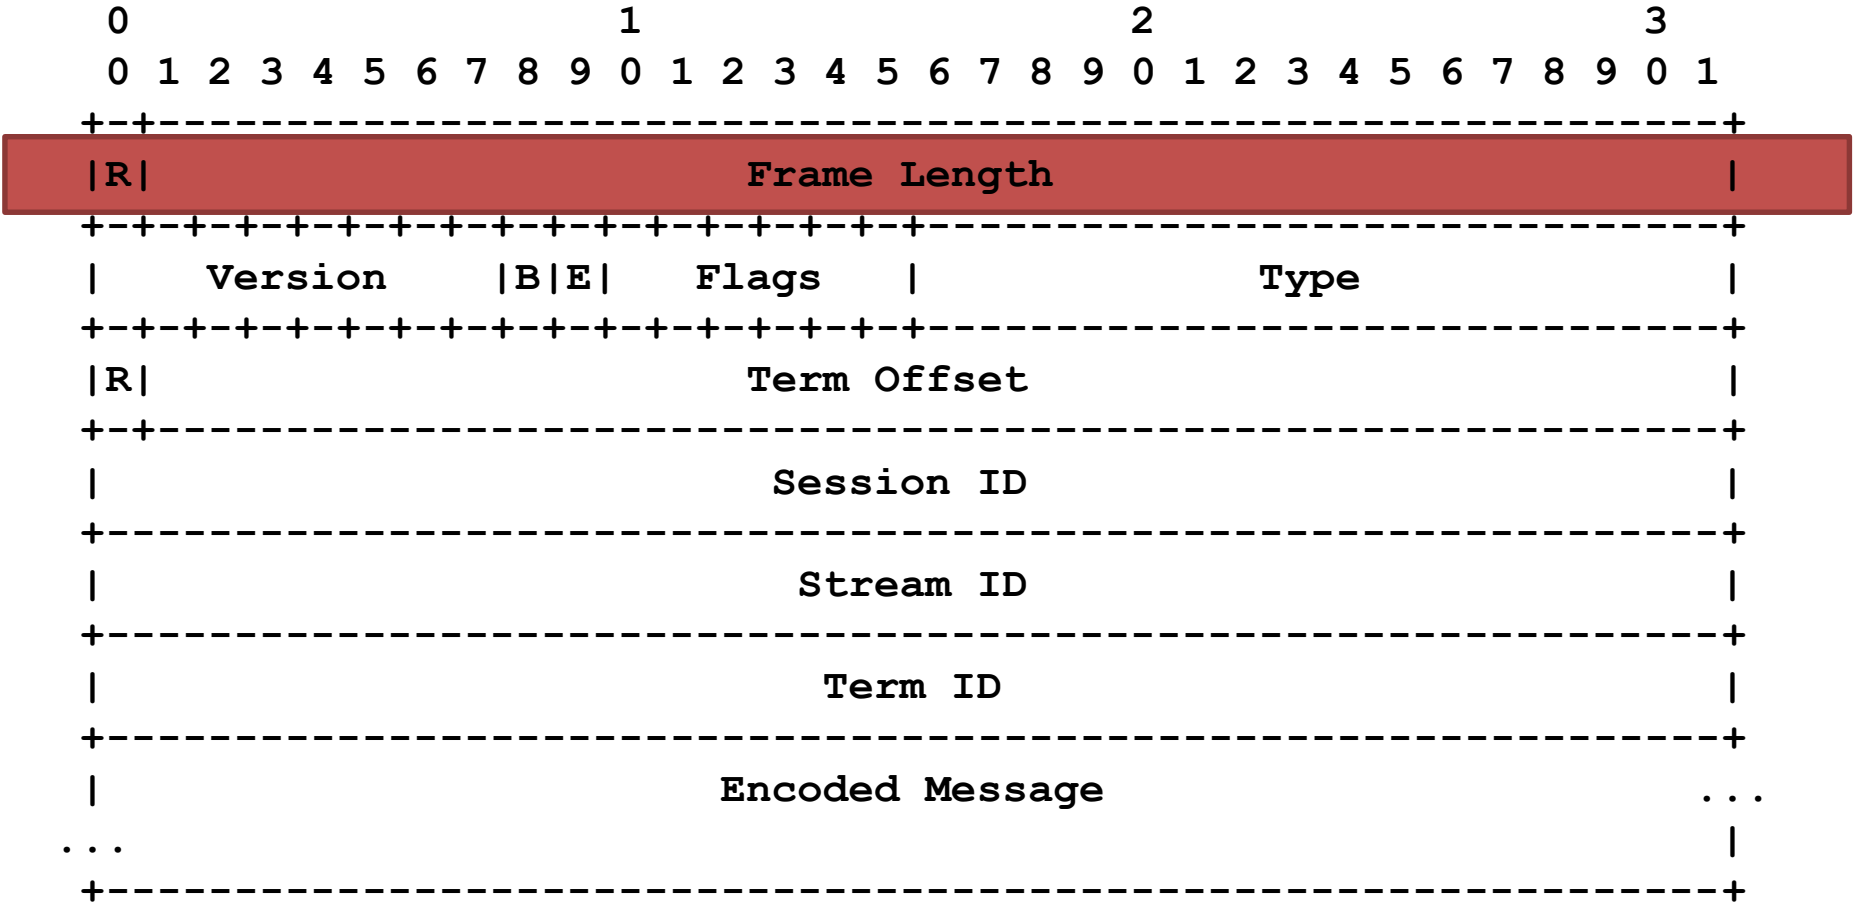


Data Message Header




How to remove a failure?

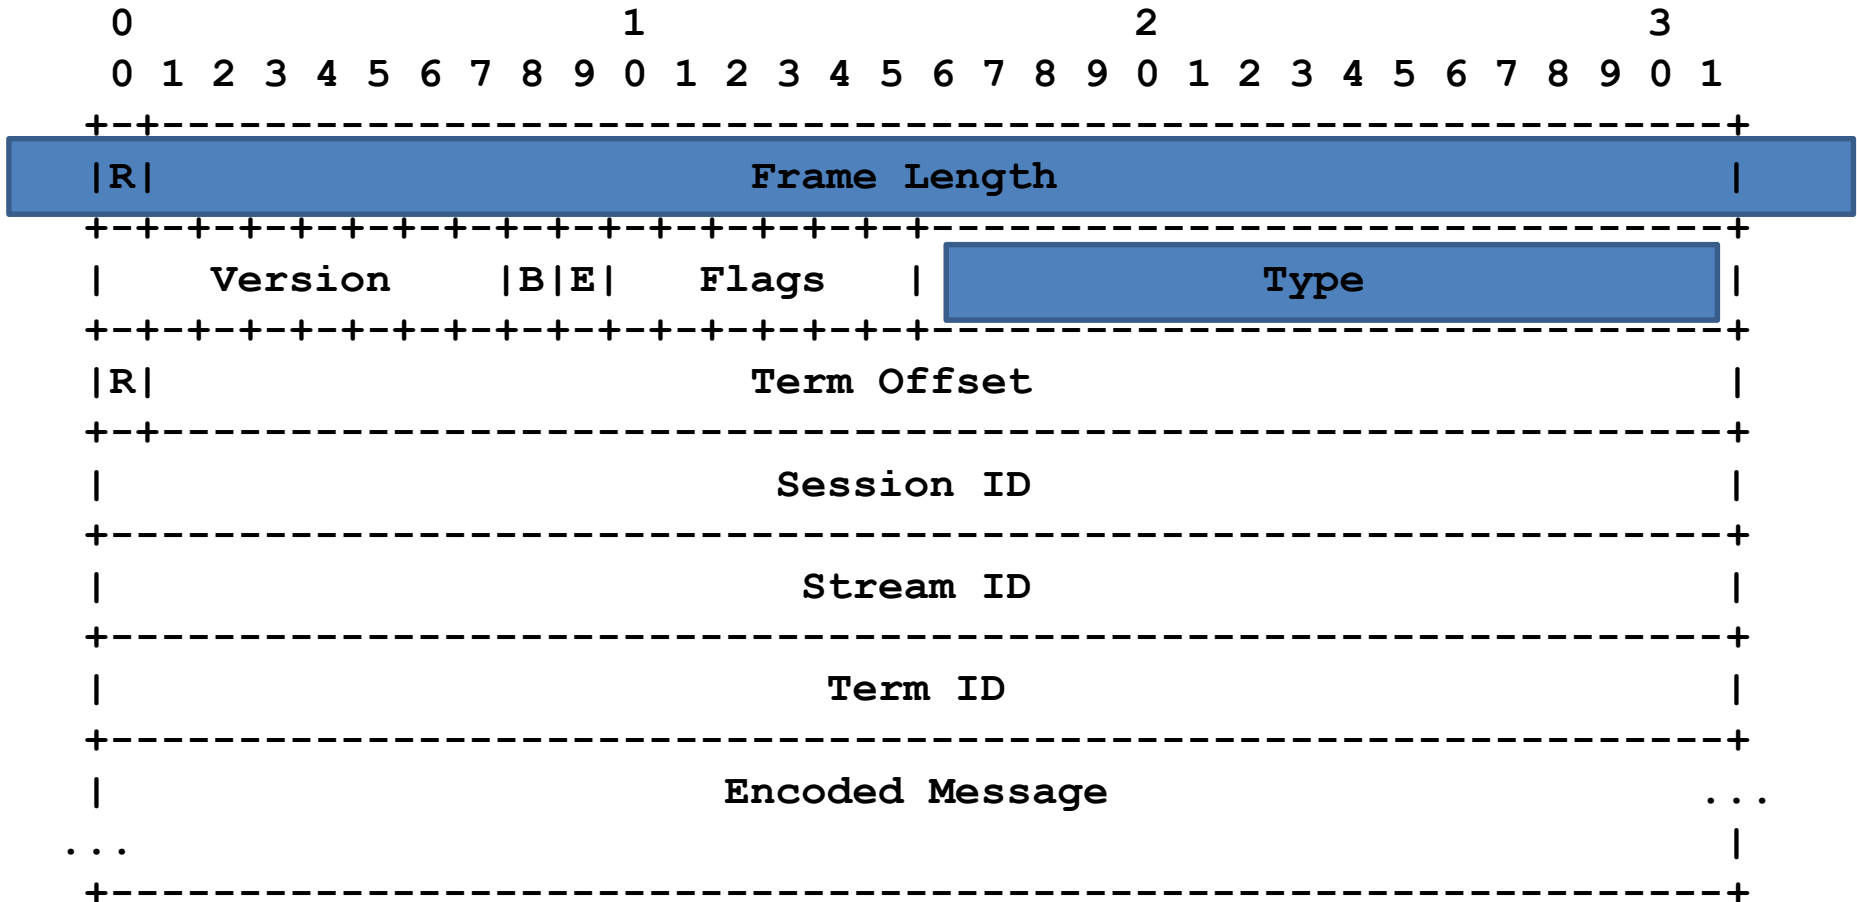
Data Message Header



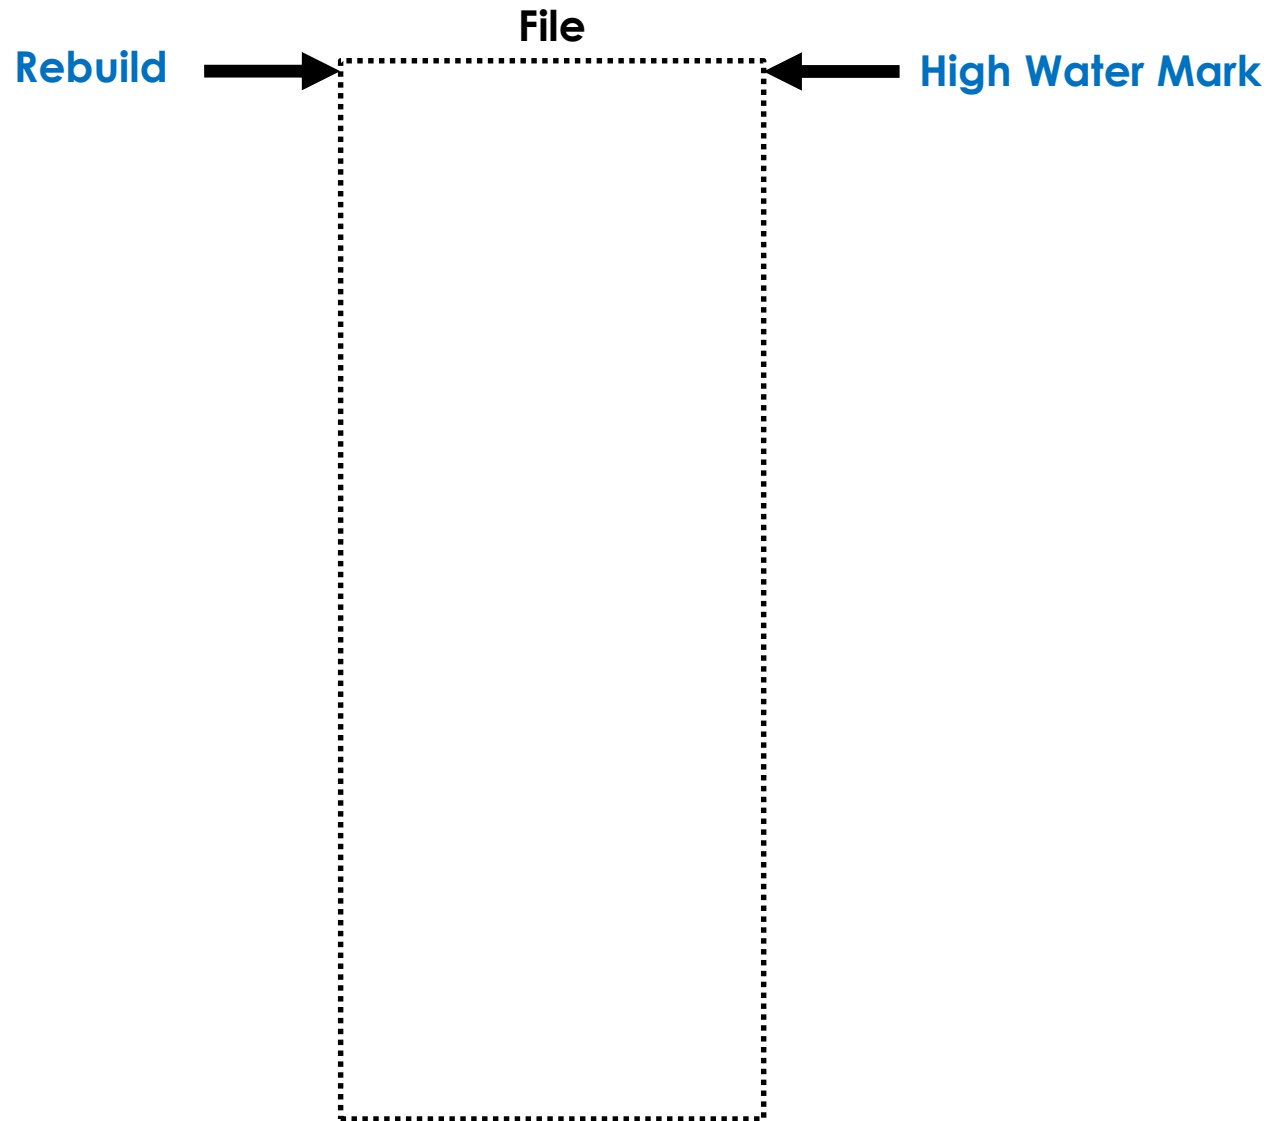
Data Message Header

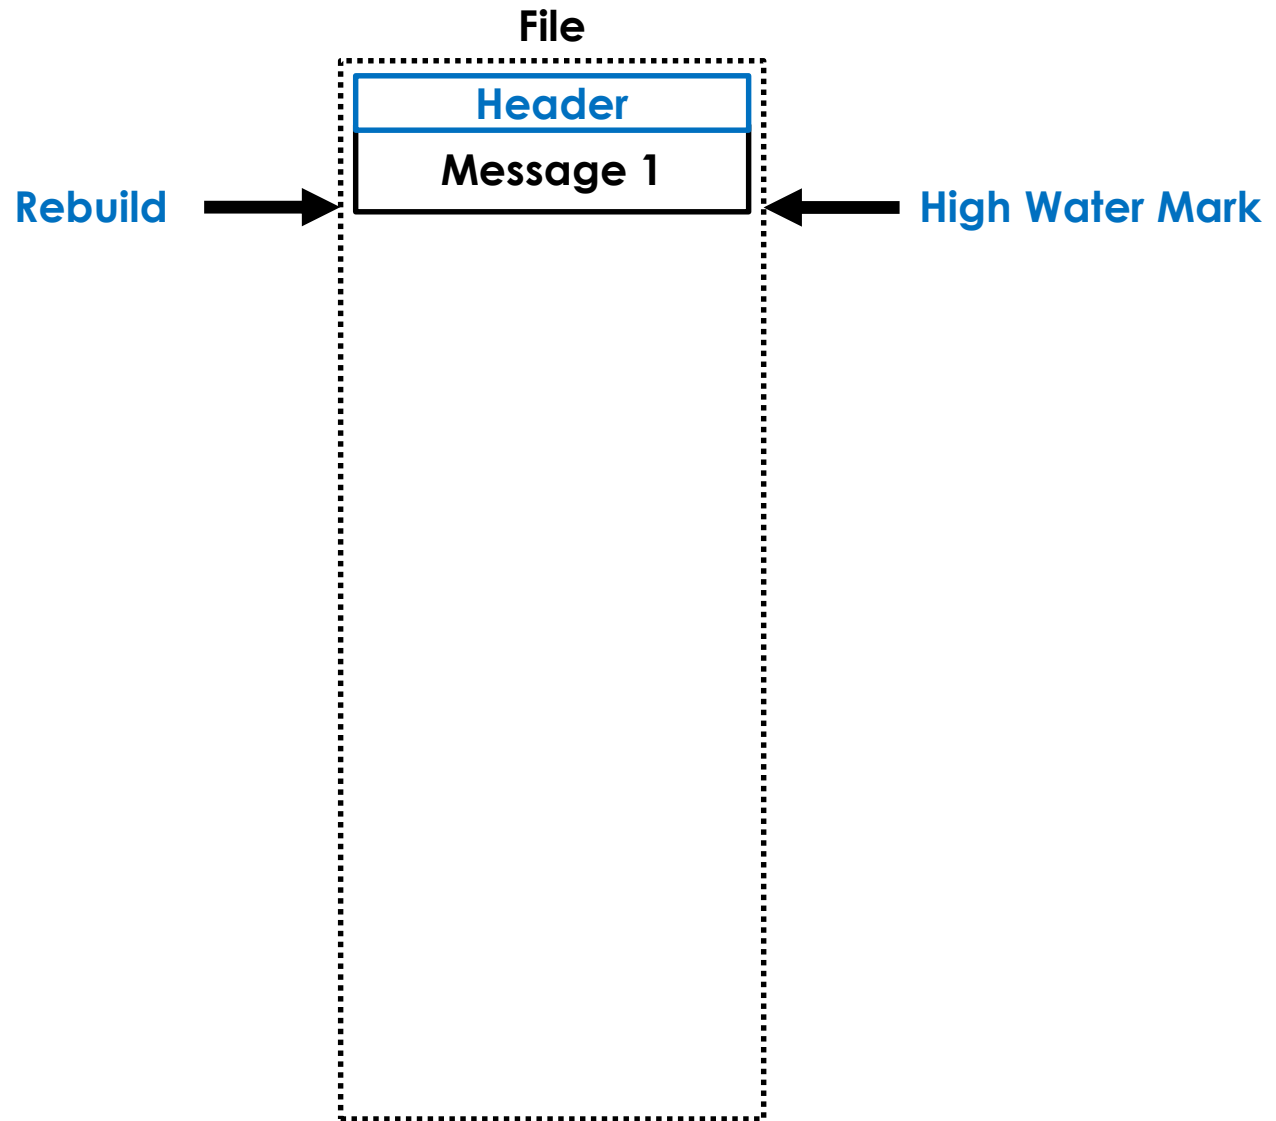


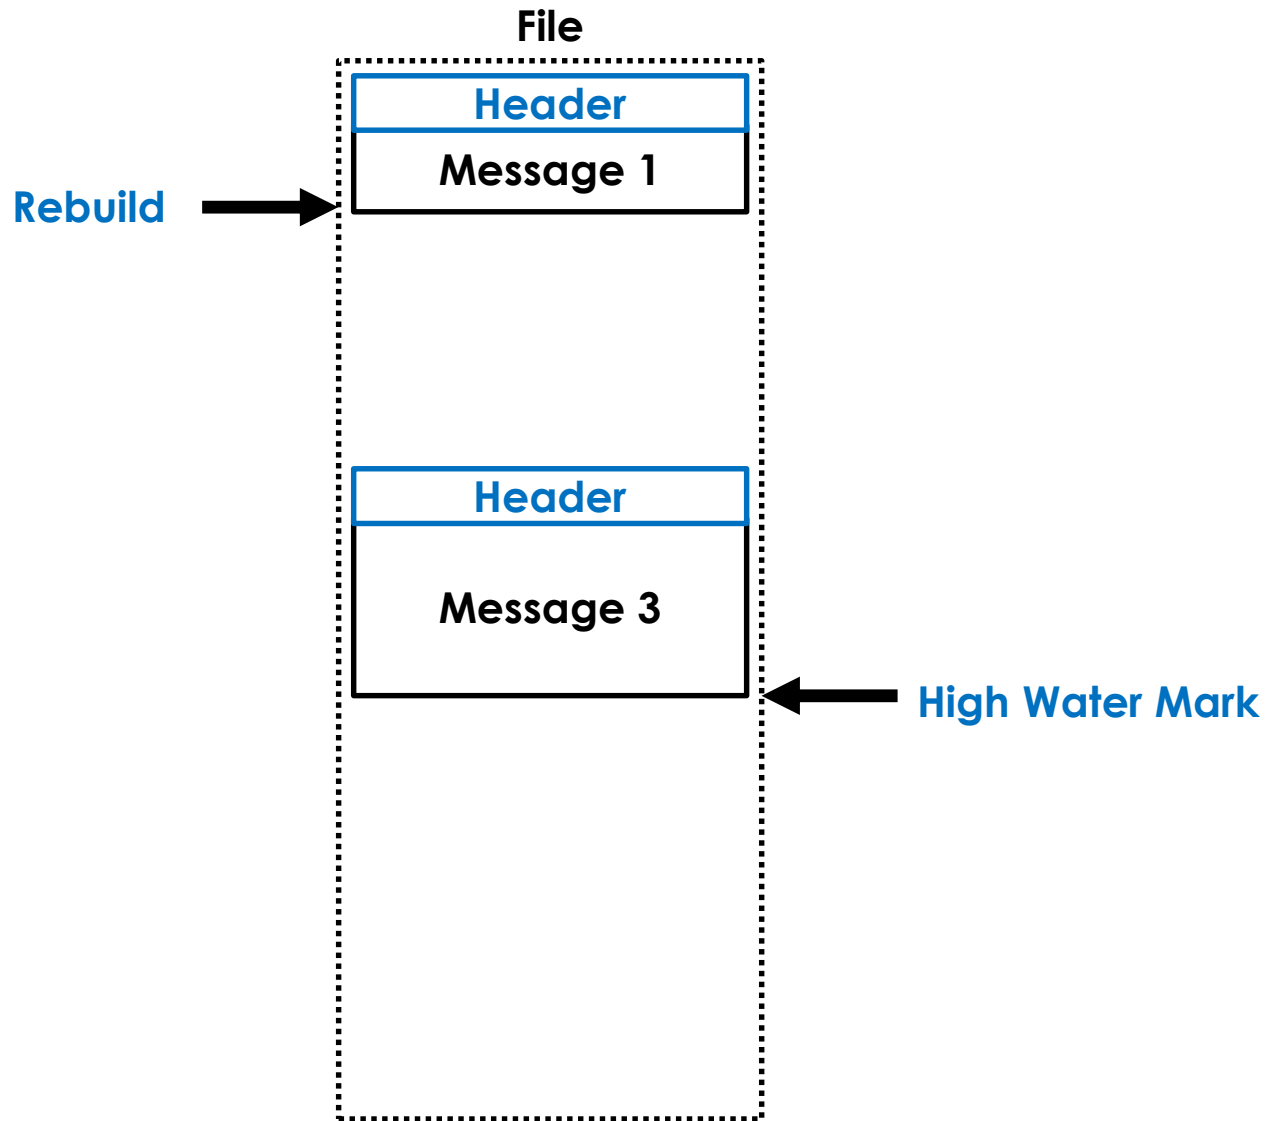
Data Message Header




How do we replicate a log?










What if a gap is not filled?

***The Log Buffer has
Strong Eventual Consistency***



***Consider the
implications of this
memory layout***

***How do we know reason about
the state of the system?***

*Publishers, Senders,
Receivers, and Subscribers
all keep position counters
which are monotonic*

In closing...

Shared Mutable State Is Evil ???

Single Writer

Shared Nothing

***Build Protocols on
Monotonically Increasing State***

<https://github.com/real-logic/aeron>

Questions?

Twitter: @mjpt777

“Travel is fatal to prejudice, bigotry, and narrow-mindedness, and many of our people need it sorely on these accounts. Broad, wholesome, charitable views of men and things cannot be acquired by vegetating in one little corner of the earth all one's lifetime.”

- Mark Twain